

AssignmentDate	
StudentName	DEVALLA CHARAN SRI SAI
StudentRollNumber	111519104022
MaximumMarks	2 Marks

#ProblemStatement:CustomerSegmentationAnalysis

###Description: You own the mall and want to understand the customers who can quickly converge [Target Customers] so that the insight can be given to the marketing team and plan the strategy accordingly.

Download and load Dataset

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import matplotlib.pyplot as plt
```

```
df = pd.read_csv('Mall_Customers.csv') #No Target Column - Unsupervised Machine Learning
df.head()
```

```
CustomerID Gender Age Annual Income (k$) Spending Score (1-100)
0 Male 19 15 39 15 81
1 2 Male 21 16 6
2 3 Female 20 16 77
3 4 Female 23 17 40
4 5 Female 31
```

```
df = df.rename(columns={'Annual Income (k$)': 'Annual_Income', 'Spending Score (1-100)': 'Spending_Score'})
df.head()
```

```
CustomerID Gender Age Annual_Income Spending_Score
0 Male 19 15 39 15 81
1 2 Male 21 16 6
2 3 Female 20 16 77
3 4 Female 23 17 40
4 5 Female 31
```

```
df.shape
```

```
(200, 5)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
```

```
# Column Non-Null Count Dtype
---
0 CustomerID 200 non-null int64
1 Gender 200 non-null object
```

```

2   Age                200non-null      int64
3   Annual_Income      200non-null      int64
4   Spending_Score     200non-null      int64
dtypes: int64(4),
object(1)memoryusage:7.9+ KB

df.Gender.unique()
array(['Male', 'Female'], dtype=object)df.Age.unique()
array([19,54,
      29, 21, 20, 23, 31, 22, 35, 64, 30, 67, 58, 24, 37, 52, 25, 46,
      51, 69, 45, 40, 60, 53, 18, 49, 42, 36, 65, 48, 50, 27, 33, 59, 47,
      41])
      70, 63, 43, 68, 32, 26, 57, 38, 55, 34, 66, 39, 44, 28, 56,
df.Gender.value_counts()Female
      112
Male      88
Name: Gender, dtype:int64

```

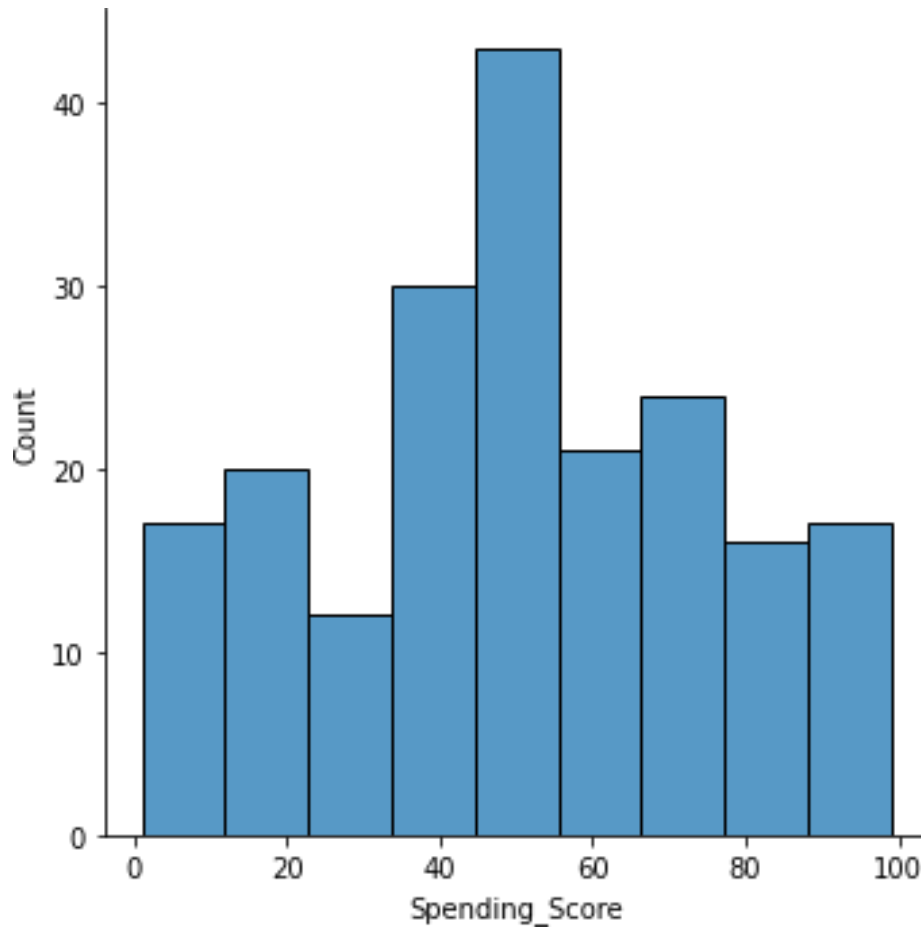
Visualizations

UnivariateAnalysis

```

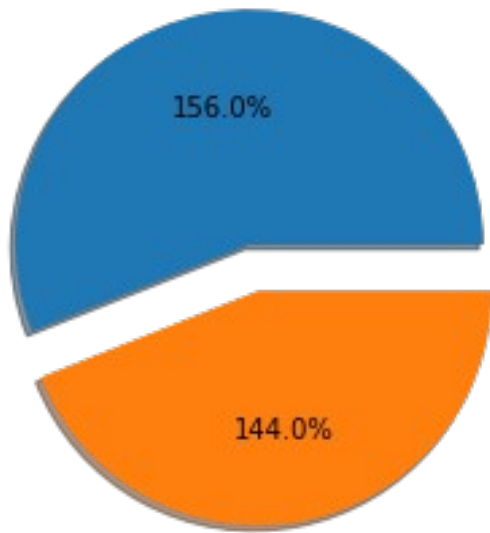
sns.displot(df.Spending_Score)
<seaborn.axisgrid.FacetGridat0x7f700626b950>

```



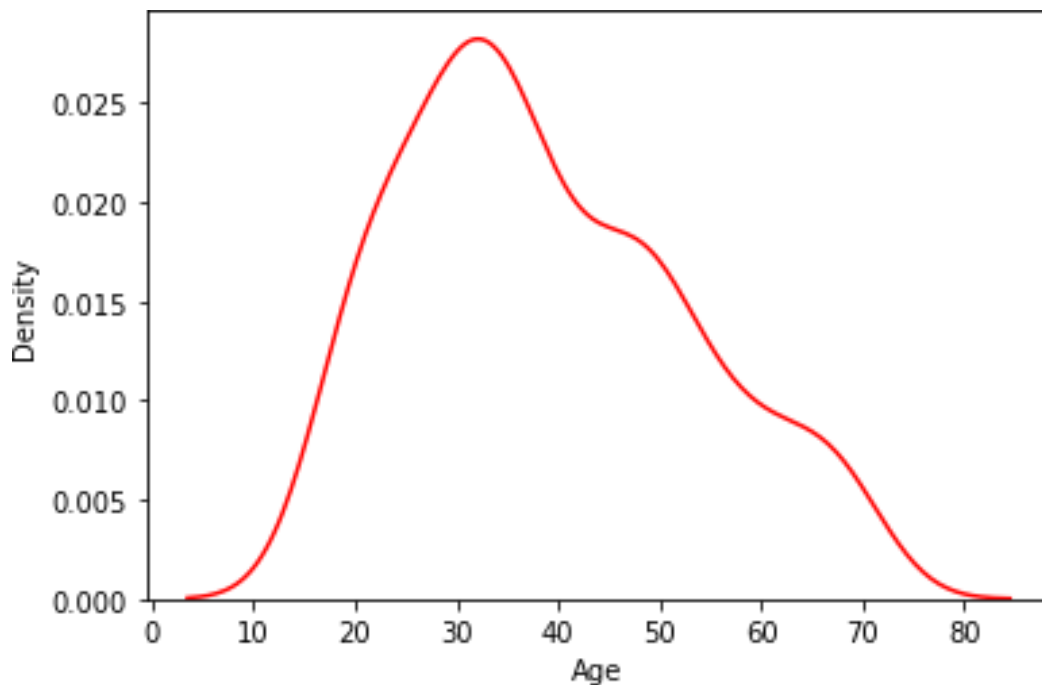
```
plt.pie(df.Gender.value_counts(),[0,0.2],shadow='True',autopct="1%.1f%
%)      #categoricalcolumn
```

```
([<matplotlib.patches.Wedge at 0x7f7005485ed0>,
 <matplotlib.patches.Wedge at 0x7f7005492950>],[Text(-
0.20611945413751356,1.080515974257694,"),
 Text(0.24359571852615253,-1.2769734241227293, ")",
 [Text(-0.11242879316591647,0.5893723495951058,'156.0%'),
 Text(0.14990505755455538,-0.7858297994601411,'144.0%')])
```



```
sns.kdeplot(df.Age,color="red")
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f700549a450>
```



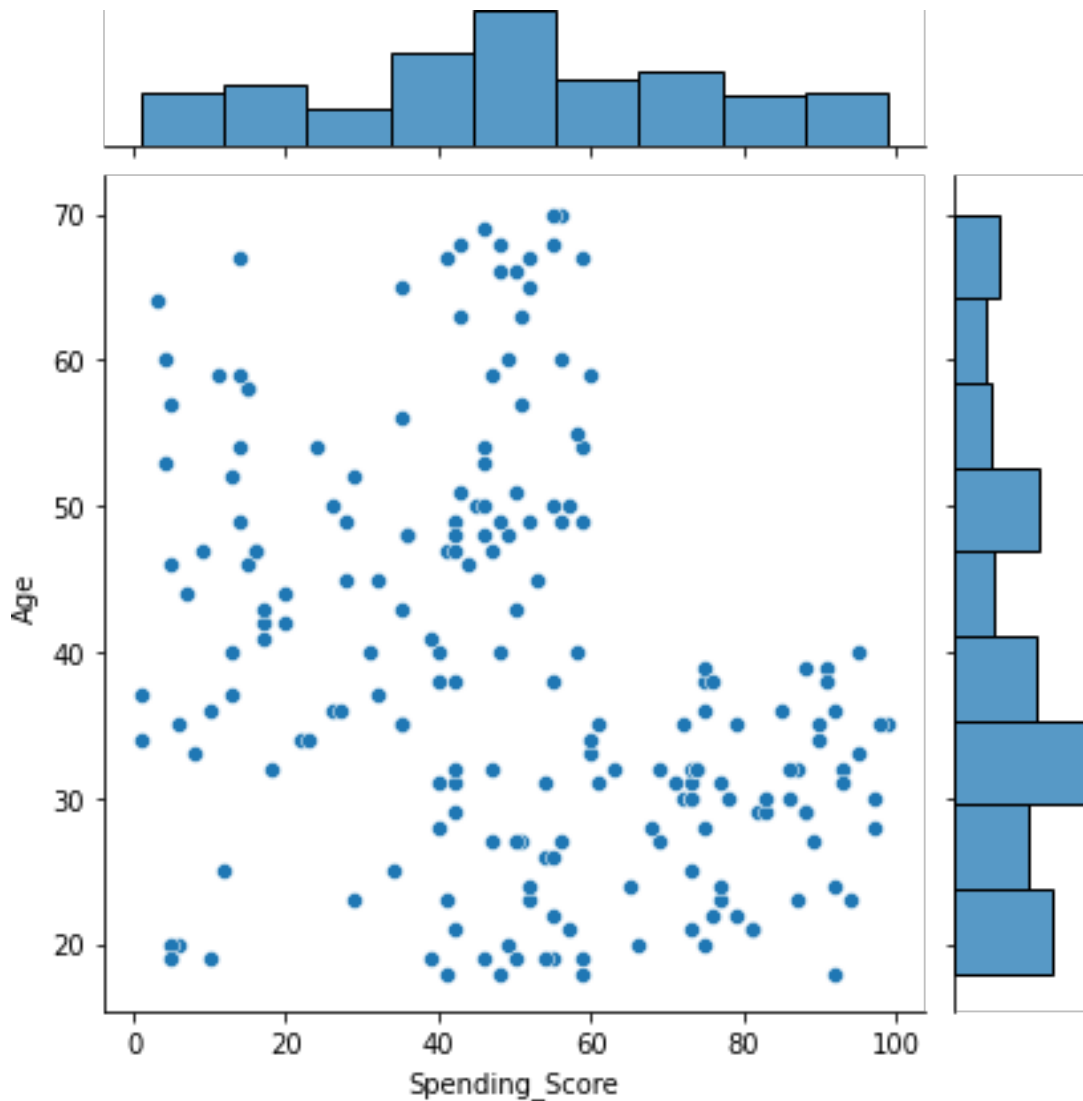
Bi-variateAnalysis

```
sns.jointplot(df.Spending_Score,df.Age)
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43:FutureWarning:Pass the following variables as keyword args:x,y.From version 0.12,the only valid positional argument will be `data`,

and passing other arguments without an explicit keyword will result in an error or misinterpretation.
FutureWarning


<seaborn.axisgrid.JointGrid at 0x7f7005459c50>



```
sns.scatterplot(df.Age,df.Annual_Income,color="green")
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43:FutureWarning:Pass the following variables as keyword args: x,y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
FutureWarning

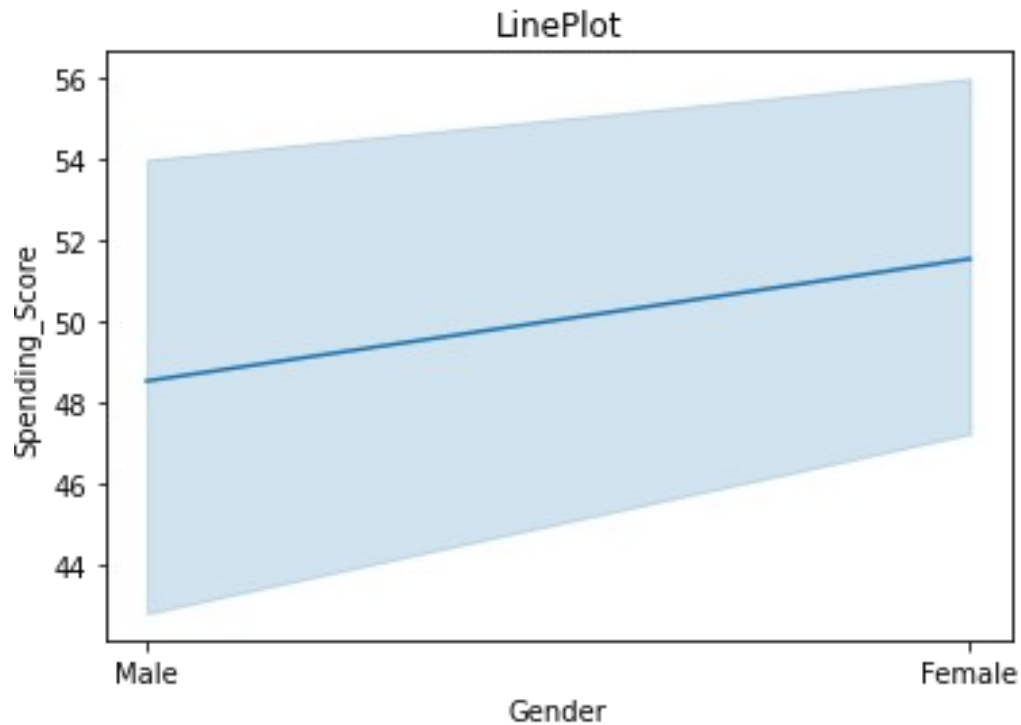
<matplotlib.axes._subplots.AxesSubplot at 0x7f7005268410>



```
sns.lineplot(df.Gender,df.Spending_Score)plt.xlabel('Gender')plt.ylabel('Spending_Score')plt.title('LinePlot')
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43:FutureWarning:Pass the following variables as keyword args: x,y.From version 0.12,the only valid positional argument will be `data`,and passing other arguments without an explicit keyword will result in an error or misinterpretation.
FutureWarning

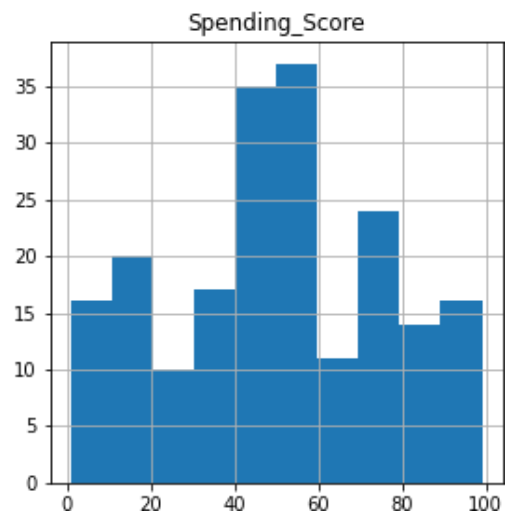
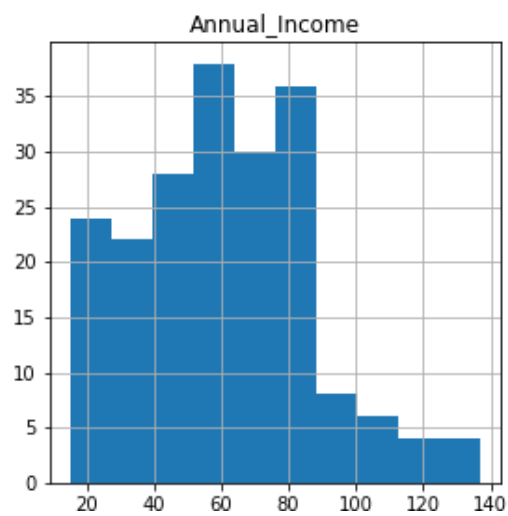
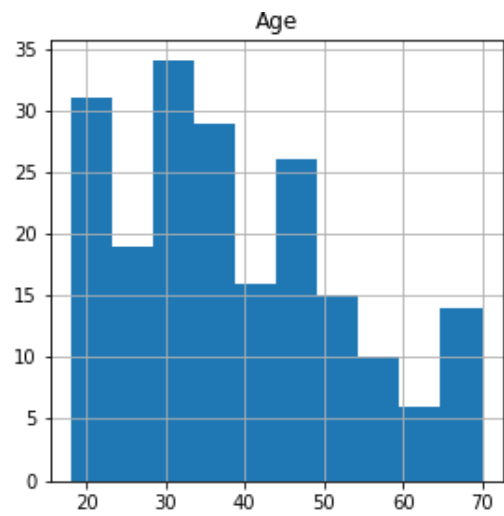
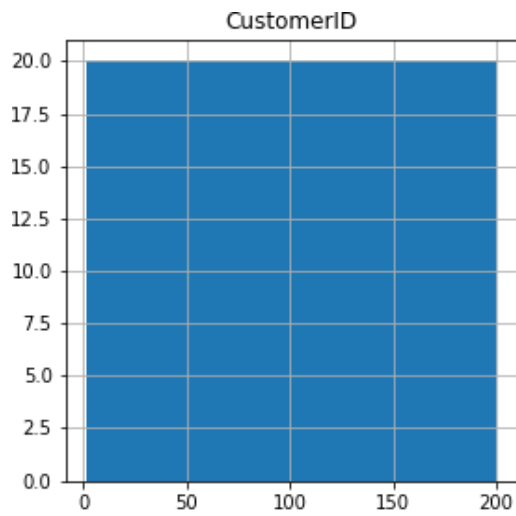
```
Text(0.5, 1.0,'LinePlot')
```



Multi-variateAnalysis

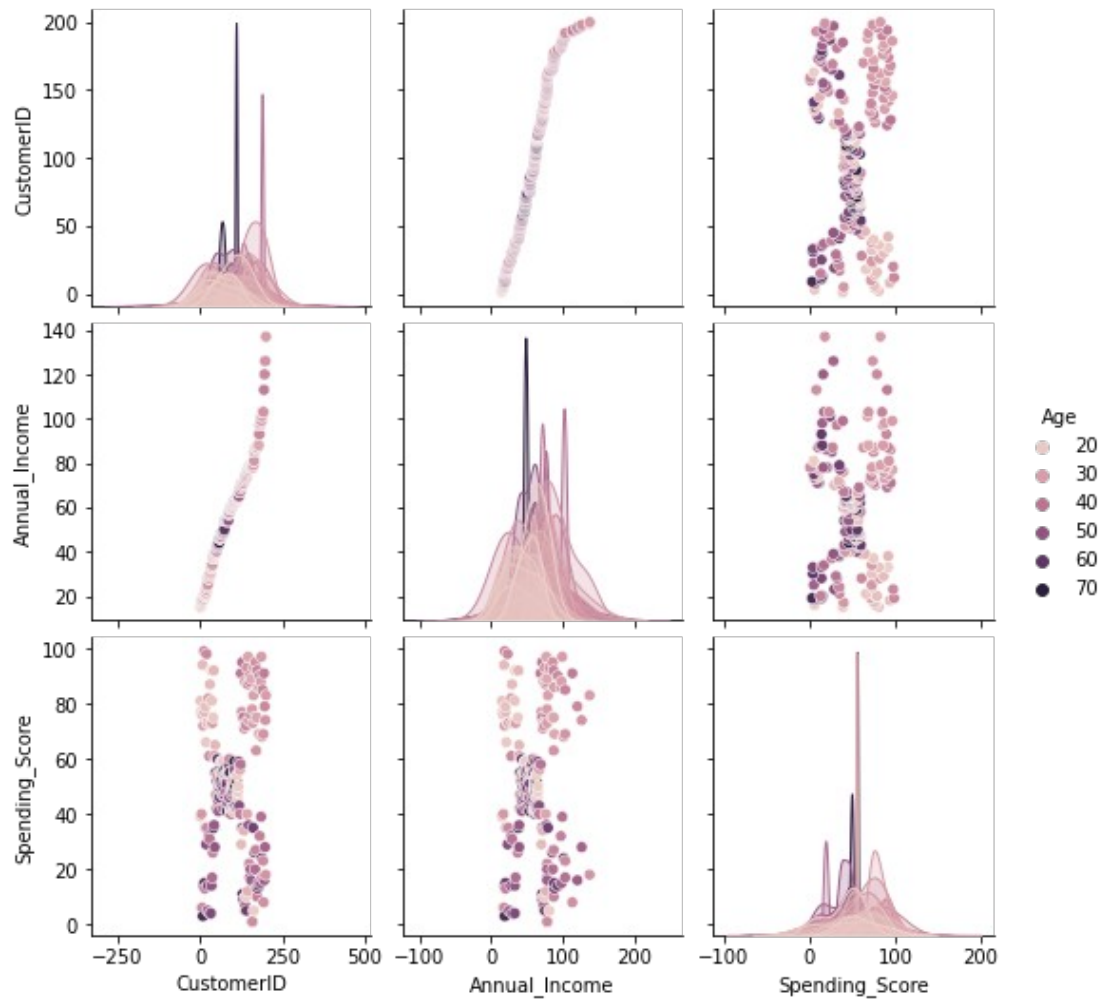
```
df.hist(figsize=(10,10))
```

```
array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f7005203910>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f70051db810>],  
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7f7005191c90>,  
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f70051541d0>]],  
      dtype=object)
```



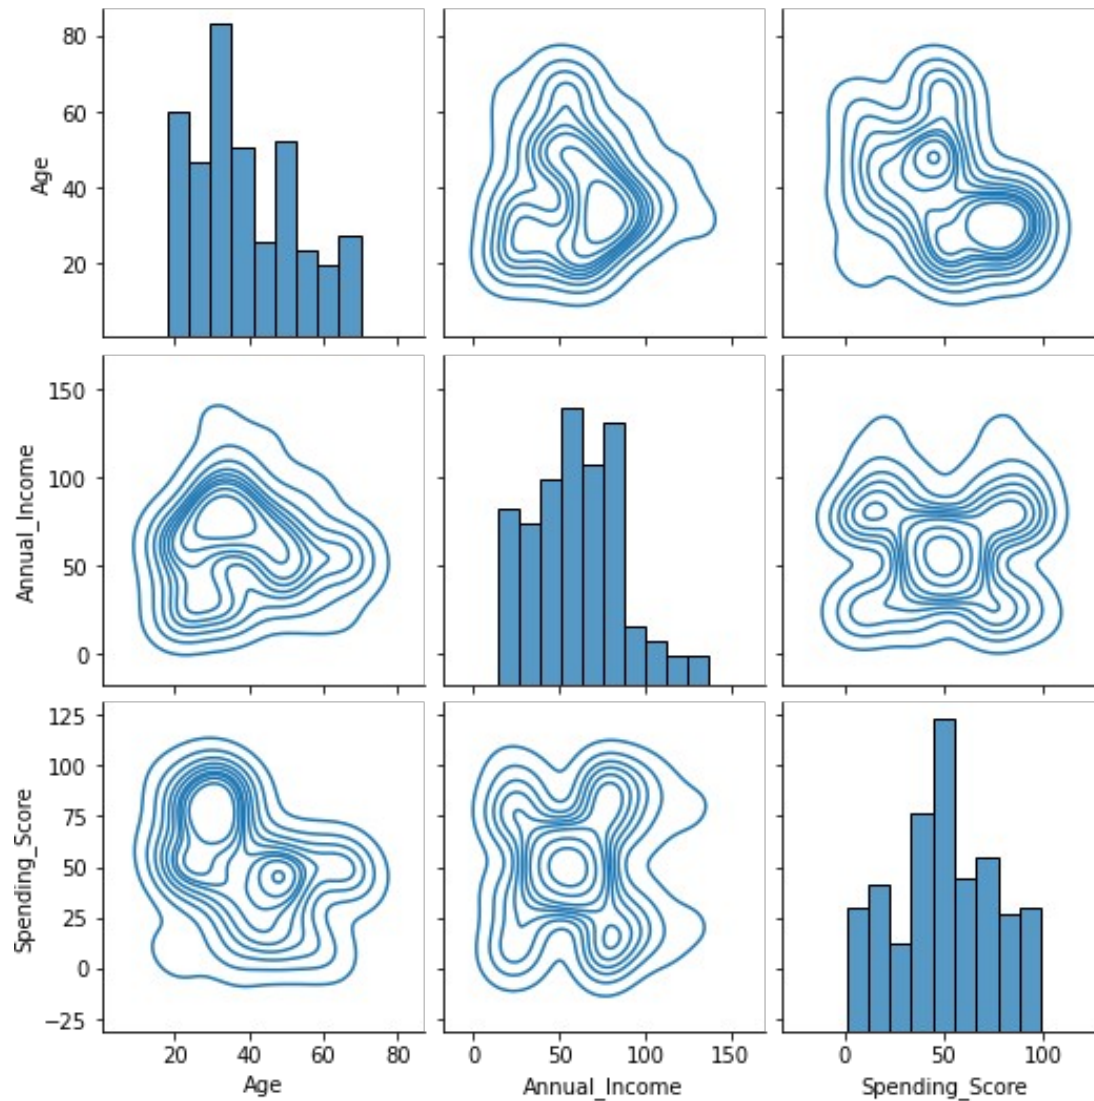
```
sns.pairplot(df,kind='scatter',hue='Age')
```

```
<seaborn.axisgrid.PairGrid at 0x7f700510cd90>
```

```
sns.pairplot(data=df[['Age','Annual_Income','Spending_Score']],kind='kde',diag_kind='hist')
```

```
<seaborn.axisgrid.PairGrid at 0x7f7004bd3cd0>
```



Descriptivestatistics

df.describe()

	CustomerID	Age	Annual_Income	Spending_Score
count	200.000000	200.000000	200.000000	200.000000
mean	100.500000	38.850000	60.560000	50.200000
std	57.879185	13.969007	26.264721	25.823522
min	1.000000	18.000000	15.000000	1.000000
25%	50.750000	28.750000	41.500000	34.750000
50%	100.500000	36.000000	61.500000	50.000000
75%	150.250000	49.000000	78.000000	73.000000
max	200.000000	70.000000	137.000000	99.000000

Handle missing data

```
df.isnull().any() #no missing data
```

```
CustomerID      False
Gender           False
Age              False
Annual_Income   False
Spending_Score  False
dtype: bool
```

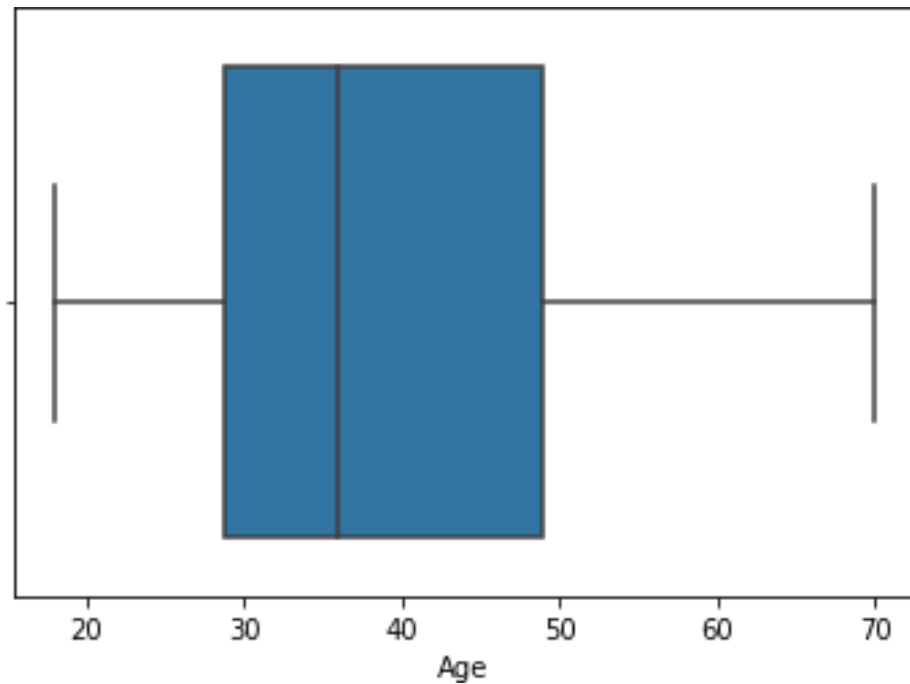
#Outliers Replacement

```
sns.boxplot(df.Age) #no outliers
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43:FutureWarning:Pass the following variables as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f7004604090>
```



Check for Categorical column and perform encoding

```
from sklearn.preprocessing import LabelEncoder
```

```
le = LabelEncoder()
```

```
df.Gender=le.fit_transform(df.Gender)df.head()
```

	CustomerID	Gender	Age	Annual_Income	Spending_Score	
	1	1	19	15	39	1
1	2	1	21	15	15	81
2	3	0	20	16	16	6
3	4	0	23	16	16	77
4	5	0	31	17	17	40

Perform clustering algorithm

```
from sklearn import cluster
```

```
error=[]
```


```
for i in range(1,11): kmeans=cluster.KMeans(n_clusters=i,init='k-means++',random_state=0) kmeans.fit(df) error.append(kmeans.inertia_)
```

```
error[9]75512.060000000
```

```
3,
```

```
387065.71377137717,  
271384.508782868,  
195401.19855991466,  
157157.7579059829,  
122625.19813553878,  
103233.01724386725,  
86053.67444777445,  
76938.97565600359,  
69231.33607611558]
```

```
import matplotlib.pyplot as plt  
plt.plot(range(1,11),error)  
plt.title('Elbow method')  
plt.xlabel('no of cluster')  
plt.ylabel('error')  
plt.show()
```



```
k_means_model=cluster.KMeans(n_clusters=3,init='k-means+
+',random_state=0)k_means_model.fit(df)KMeans(n
_clusters=3,random_state=0)
clustered_data=k_means_model.predict(df)
```

Add the cluster data with the primary dataset

```
df['Clustered_data'] = pd.Series(clustered_data)df.head()
```

	CustomerID	Gender	Age	Annual_Income	Spending_Score
Clustered_data					
0	1	1	19	15	39
0					
1	2	1	21	15	81
0					
2	3	0	20	16	6
0					
3	4	0	23	16	77
0					
4	5	0	31	17	40
0					

Split the data into dependent and independent variables

```
y=df['Clustered_data']
```

y *#y - target columns*

```
0    0
1    0
2    0
3    0
4    0
```

```
..
195   2
196   2
197   2
198   2
199   2
```

Name: Clustered_data, Length: 200, dtype: int32

```
X=df.drop(columns=['Clustered_data'],axis=1)
```

X.head() *#X-predicting columns*

	CustomerID	Gender	Age	Annual_Income	Spending_Score	0	1
	1	1	19	15	39		
1	2	1	21	15		81	
2	3	0	20	16		6	
3	4	0	23	16		77	
4	5	0	31	17		40	

Scale the independent variables

```
from sklearn.preprocessing import scale
```

```
data=pd.DataFrame(scale(X),columns=X.columns)data.head()
```

	CustomerID	Gender	Age	Annual_Income	Spending_Score
0	-1.723412	1.128152	-1.424569	-1.738999	-0.434801
1	-1.706091	1.128152	-1.281035	-1.738999	1.195704
2	-1.688771	-0.886405	-1.352802	-1.700830	-1.715913
3	-1.671450	-0.886405	-1.137502	-1.700830	1.040418
4	-1.654129	-0.886405	-0.563369	-1.662660	-0.395980

Split the data into training and testing

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(data,y,test_size=0.3,random_state=1)
```

X_train.shape,X_test.shape((140,5),

(60,5))

```
y_train.shape,y_test.shape((140,),(6
0,))
```

Buildthemodel

```
from sklearn.neighbors import
KNeighborsClassifiermodel=KNeighborsClassifier()

model.fit(X_train,y_train)          #K -Nearest Neighbourmodel (KNN)

KNeighborsClassifier()
```

Trainthemodel

```
pred_train= model.predict(X_train)pred_train
```

```
array([1,      1,  1,  0,  0,  0,  2,  1,  0,  1,  0,  1,  2,  2,  2,  1,  0,  1,  1,  1,  2,
1,
      1,  1,  2,  0,  1,  1,  2,  0,  1,  0,  2,  2,  2,  1,  2,  2,  2,  2,  1,  0,  1,
2,
1,      0,  1,  1,  2,  0,  1,  0,  2,  1,  1,  1,  2,  1,  2,  0,  1,  1,  1,  2,  2,  2,
2,      2,  2,  2,  0,  0,  1,  2,  1,  2,  0,  2,  0,  2,  1,  2,  2,  1,  2,  1,  0,  0,
1,      1,  1,  0,  0,  1,  0,  0,  0,  2,  0,  2,  1,  2,  0,  1,  1,  2,  0,  1,  2,  0,
2,      0,  1,  1,  0,  2,  2,  1,  1,  1,  0,  2,  2,  2,  2,  2,  1,  0,  2,  0,  2,  1,
2,      2,  2,  1,  2,  2,  1,  2,  0],dtype=int32)
```

Testthedata

```
y_test
```

```
58      0
40      0
34      0
102     1
184     2
198     2
95      1
4       0
29      0
168     2
171     2
18      0
11      0
89      1
```

110	1
118	1
159	2
35	0
136	2
59	0
51	0
16	0
44	0
94	1
31	0
162	2
38	0
28	0
193	2
27	0
47	0
165	2
194	2
177	2
176	2
97	1
174	2
73	1
69	1
172	2
108	1
107	1
189	2
14	0
56	0
19	0
114	1
39	0
185	2
124	1
98	1
123	1
119	1
53	0
33	0
179	2
181	2
106	1
199	2
138	2

Name: Clustered_data, dtype:int32

pred_test=model.predict(X_test)pred_test


```
array([0,0, 1, 0, 1, 2, 2, 1, 0, 0, 2, 2, 0, 0, 1, 1, 1,2,0,2,1,1,
0,
0, 1, 0, 2, 0, 0, 2, 0, 0, 2, 2, 2, 2, 1, 2, 1,0,2,1,1,2,
0,
0, 1, 0, 2, 1, 1, 1, 1, 1, 0, 2, 2, 1, 2, 2],dtype=int32)
pred=
pd.DataFrame({'Actual_value':y_test,'Predicted_value_using_KNN':pred_test})
pred.head()
```

	Actual_value	Predicted_value_using_KNN
58	0	0
40	0	1
34	0	0
102	1	1
184	2	2

Measure the performance using metrics

```
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

#Accuracy Score

```
print('Training accuracy:
',accuracy_score(y_train,pred_train))print('Testing accuracy:',accuracy_score(y_test,pred_test))
```

Training accuracy:0.9214285714285714

Testing accuracy:0.9166666666666666

#Confusion Matrix

```
pd.crosstab(y_test,pred_test)
```

col_0	Clustered_dat	0	1	2
a				
0		19	4	0
1		1	16	0
2		0	0	20

#Classification Report

```
print(classification_report(y_test,pred_test))
```

	precision	recall	f1-score	support
0	0.95	0.83	0.88	23
1	0.80	0.94	0.86	17
2	1.00	1.00	1.00	20
accuracy			0.92	60
macroavg	0.92	0.92	0.92	60

weightedavg	0.92	0.92	0.92	60
-------------	------	------	------	----